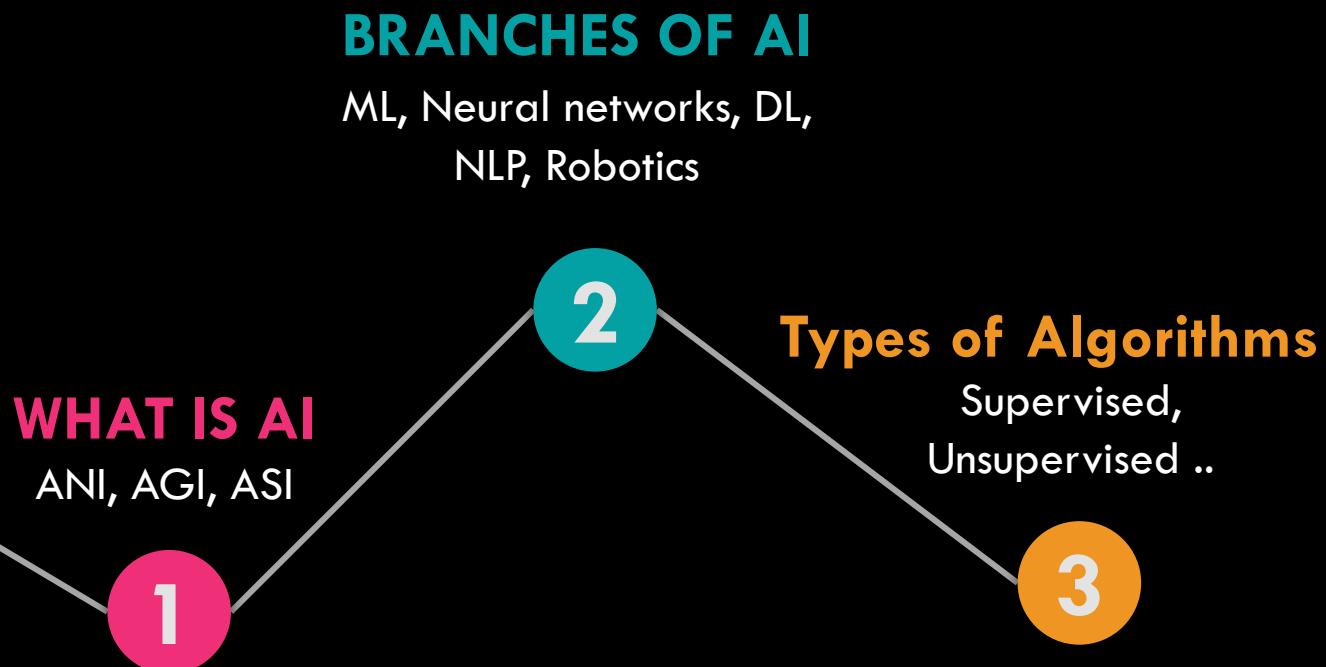


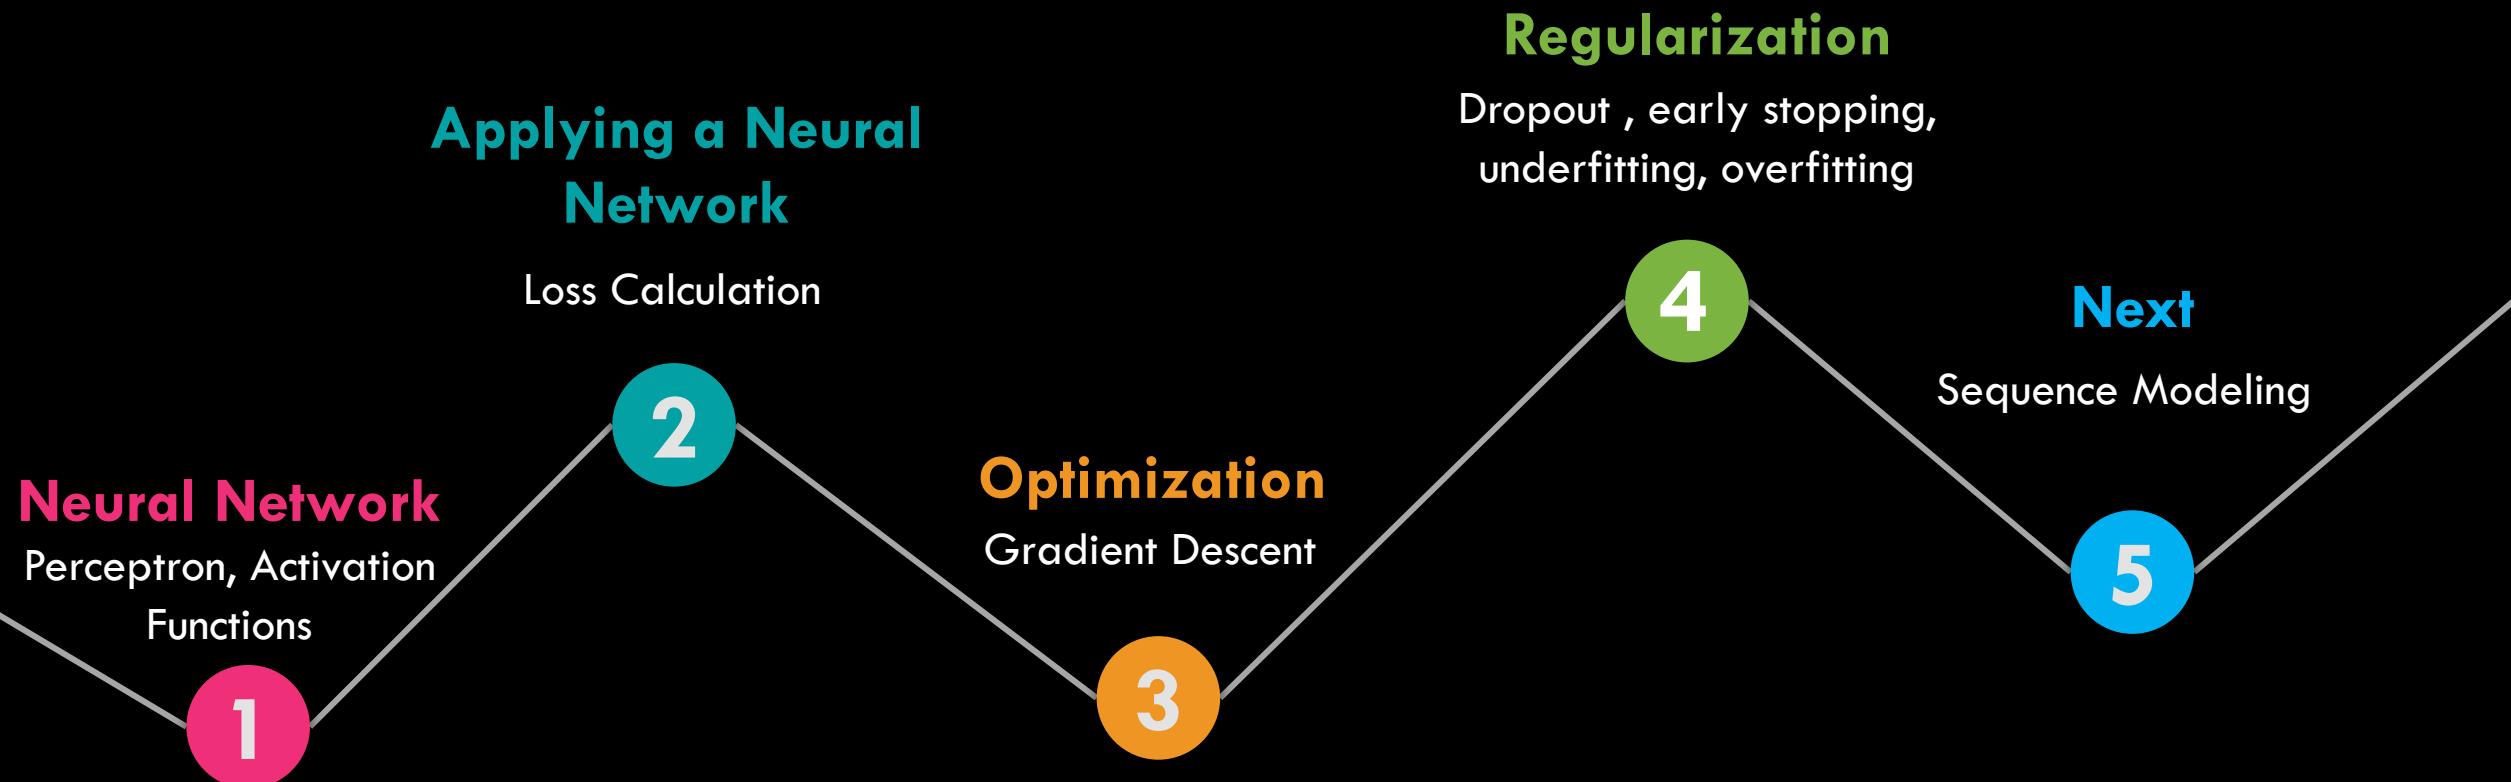
Fundamentals of AI/ML

Session-2

Recap



TODAY – Introduction to Deep Learning

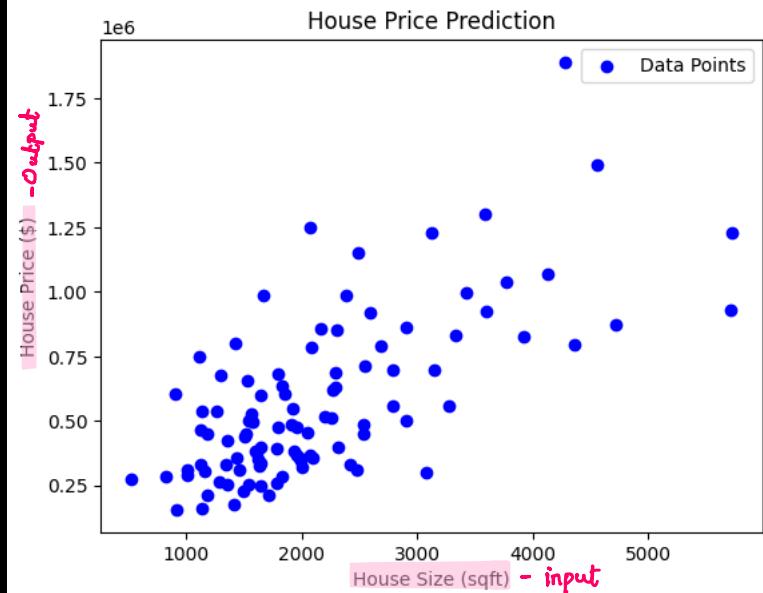


Neural Network

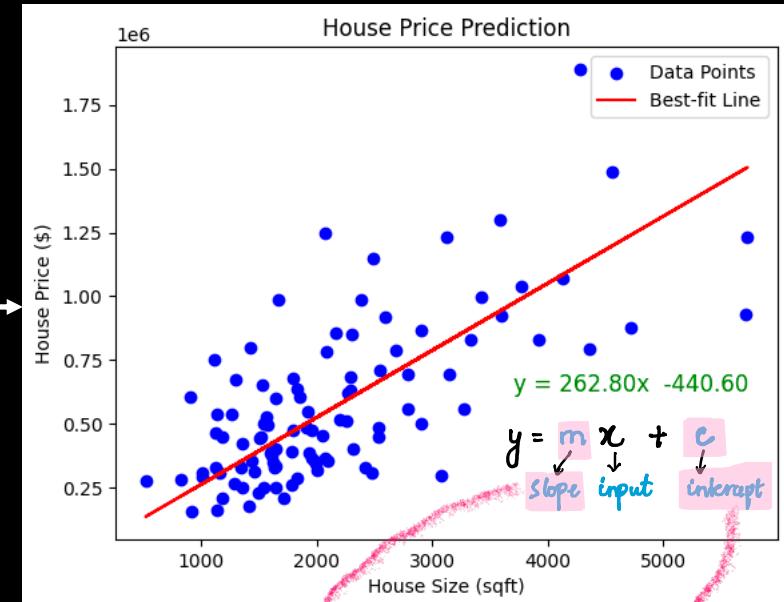


We can think of Neural Network as,
A line fitting model at its simplest.

Dataset	
sqft_living	price
2250.00	292500.00
2420.00	301000.00
3250.00	951000.00
1850.00	430000.00
2150.00	650000.00
1260.00	289000.00
2519.00	505000.00
154.00	549000.00
1660.00	425000.00
2770.00	317625.00
2720.00	975000.00
2240.00	287000.00
1000.00	204000.00
3200.00	1330000.00
4770.00	1040000.00
1260.00	325000.00
2750.00	571000.00
2380.00	360000.00
1790.00	349000.00
3430.00	832500.00

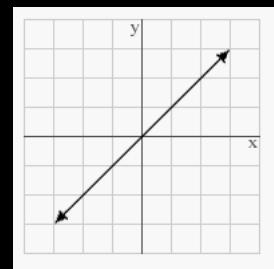


Train a **neural network** to
predict House **Price** given
House size as **Input**

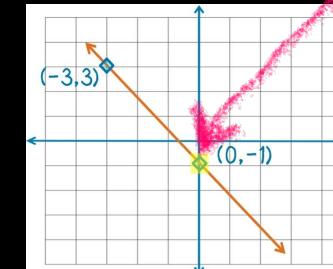


Terminology
Alert!

Slope in neural networks is weight
intercept in neural network is Bias



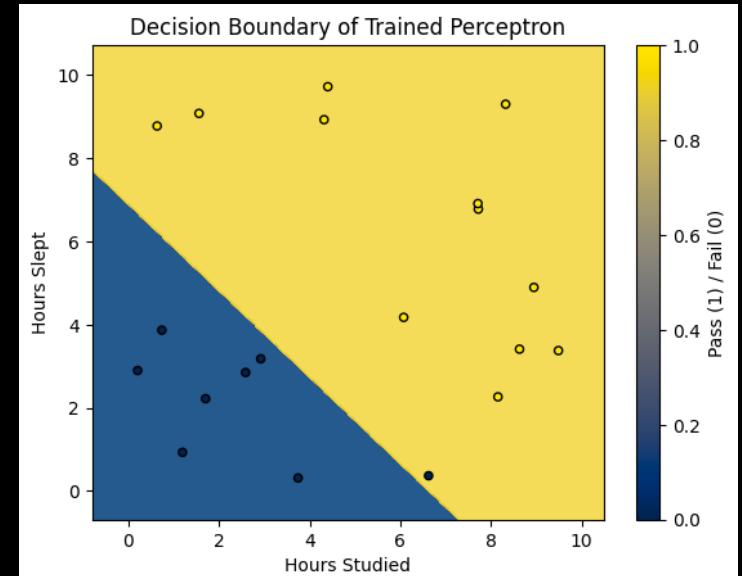
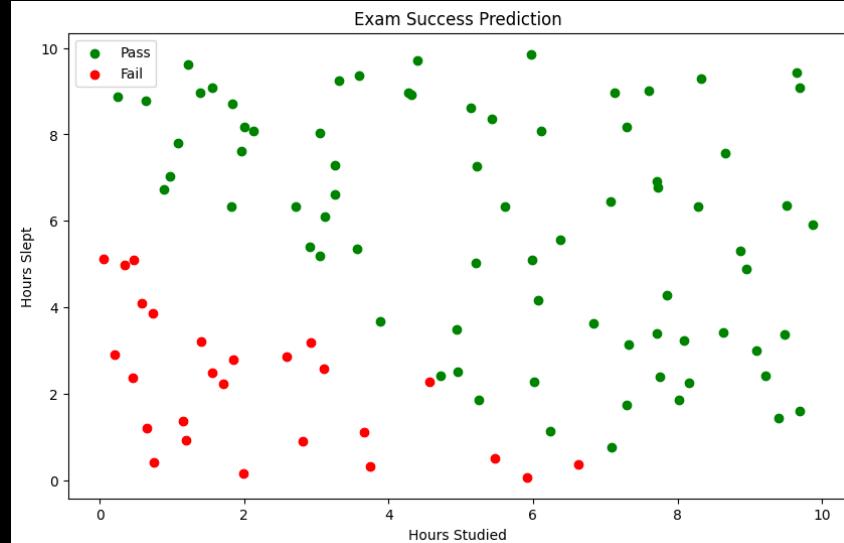
This Photo by Unknown Author is licensed under CC BY-SA



Neural Network



HoursStudied	HoursSlept	Pass
3.75	0.31	0.00
9.51	6.36	1.00
7.32	3.14	1.00
5.99	5.09	1.00
1.56	9.08	1.00
1.56	2.49	0.00
0.58	4.10	0.00
8.66	7.56	1.00
6.01	2.29	1.00
7.08	0.77	1.00
0.21	2.90	0.00
9.70	1.61	1.00
8.32	9.30	1.00
2.12	8.08	1.00
1.82	6.33	1.00
1.83	8.71	1.00
3.04	8.04	1.00
5.25	1.87	1.00
4.32	8.93	1.00
2.91	5.39	1.00



Terminology Alert!

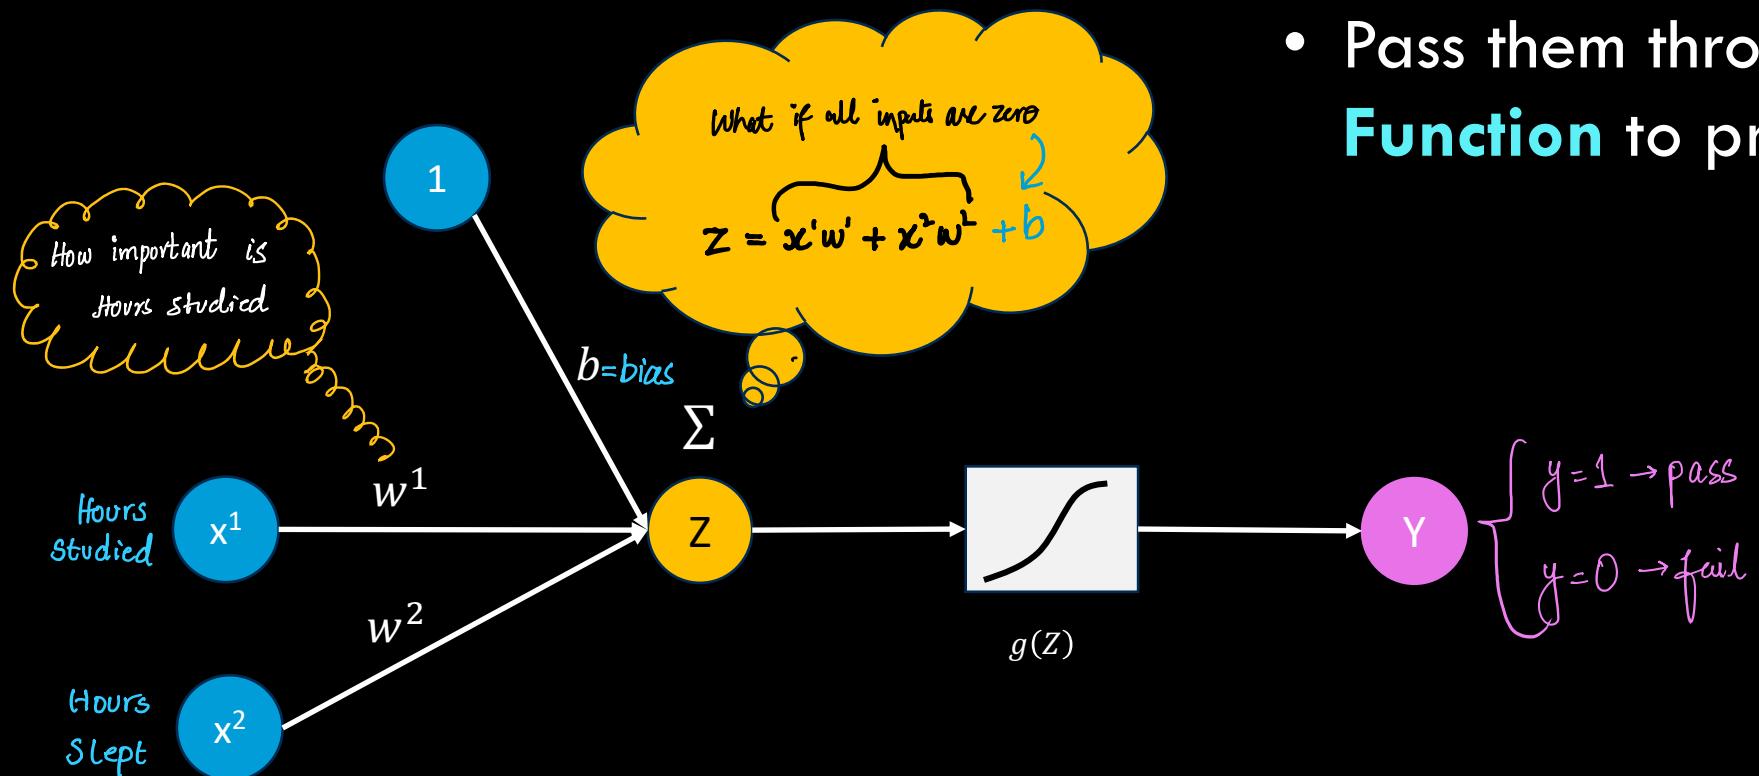
This simplest model
is referred as

Perceptron

Train a **neural network** to
predict if a student will **Pass**
given time slept and studied
as **Input**

$$\text{Pass} = g(w' \times \text{hours_studied} + w^2 \times \text{hours_slept} + b)$$

Perceptron



INPUTS

WEIGHTS
And bias

SUMMATION

Activation Function

OUTPUT

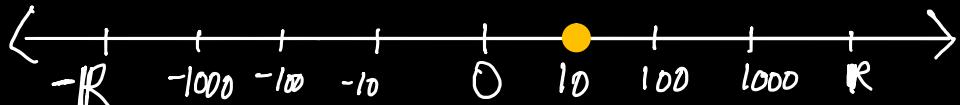
A single computational Unit of Neural Networks;

- Takes **Inputs**
- Applies **Weights and bias** to them
- Pass them through **Activation Function** to produce an **Output**

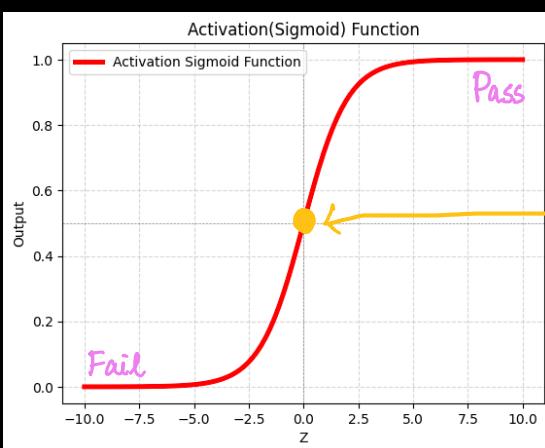
Activation Function



$$Z = x^1 w^1 + x^2 w^2 + (x^0 w^0)$$



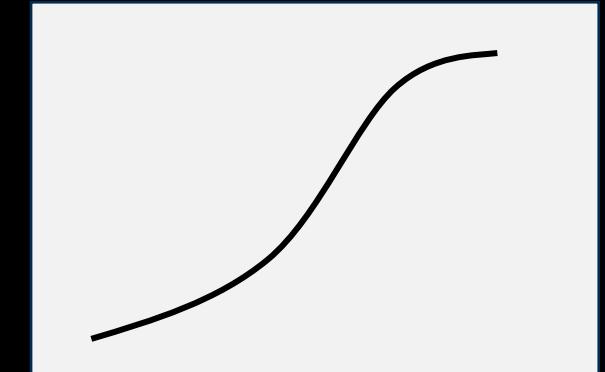
$$g(Z) = \frac{1}{1 + e^{-Z}}$$



```
import torch.nn as nn  
input_size = 2  
fc = nn.Linear(input_size, 1) # Single output neuron
```

```
x = torch.Tensor([0.6356, 8.7734]) # Hours Studied, Hours Slept  
z = fc(x)
```

```
y = torch.sigmoid(z) # g(z)  
y  
✓ 0.0s  
tensor([0.5160], grad_fn=<SigmoidBackward0>)
```



$g(Z)$



0.5 is neutral.

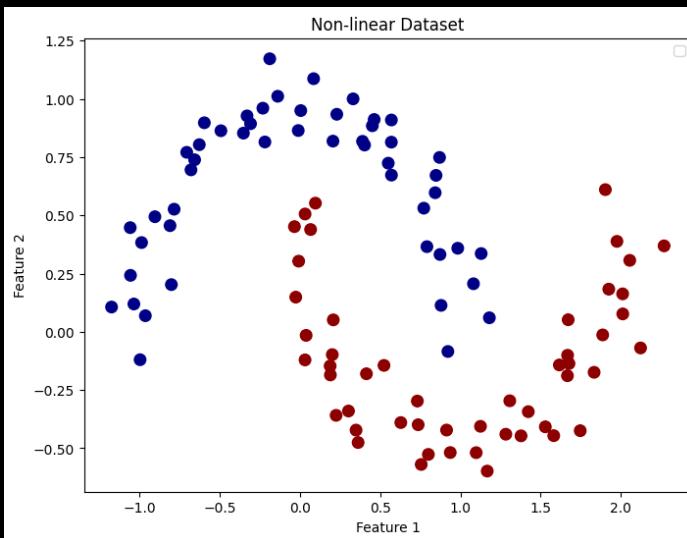
50% confident that student failed

50% confident that student Passed

Importance of Activation Function

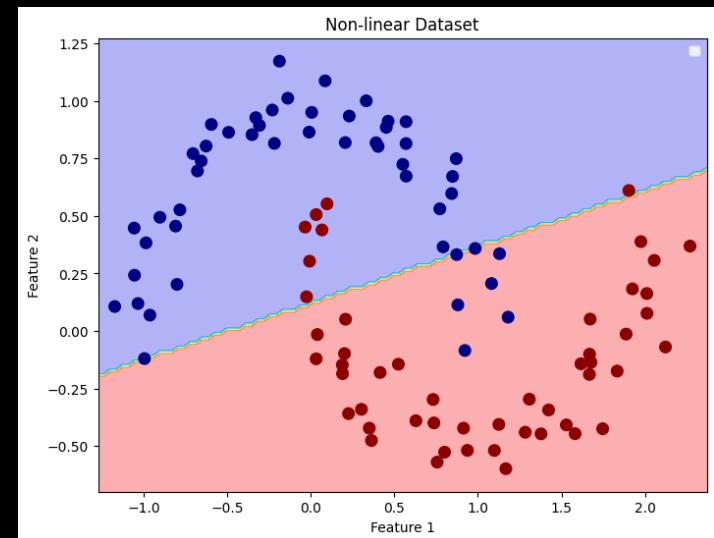


Activation Function introduces **non-linearity** to the network



```
# Define a multi-layer perceptron (MLP) model
class MLP(nn.Module):
    def __init__(self):
        super(MLP, self).__init__()
        self.fc1 = nn.Linear(2, 10) # Input 1
        self.fc2 = nn.Linear(10, 2)
        self.fc3 = nn.Linear(2, 1)

    def forward(self, x):
        x = self.fc1(x)
        x = self.fc2(x)
        x = self.fc3(x)
        return x
```



Without activation function, model will
learn linear decision boundary , no matter
the **network size**

Importance of Activation Function

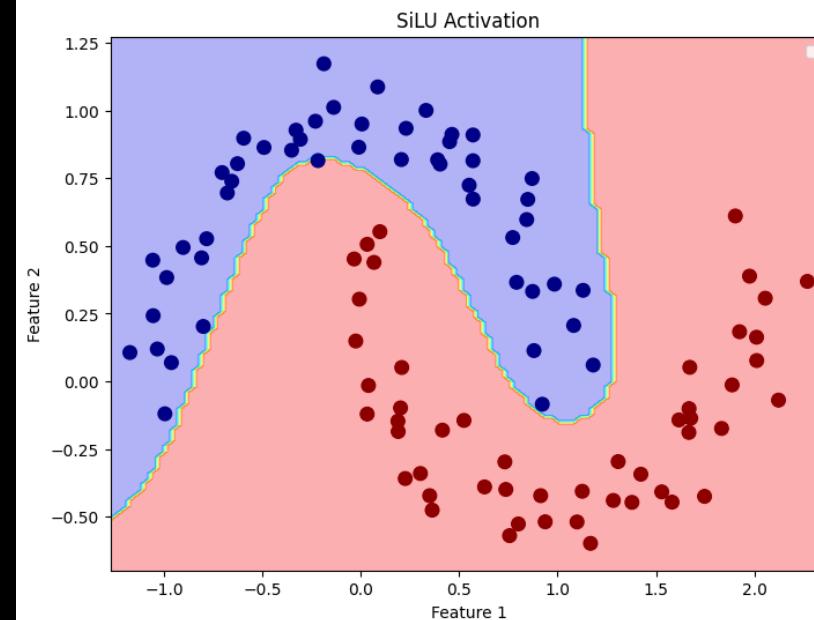
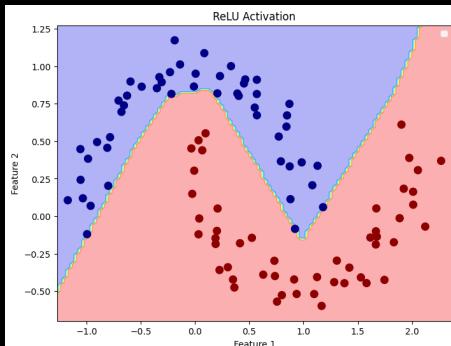


$g(Z)$

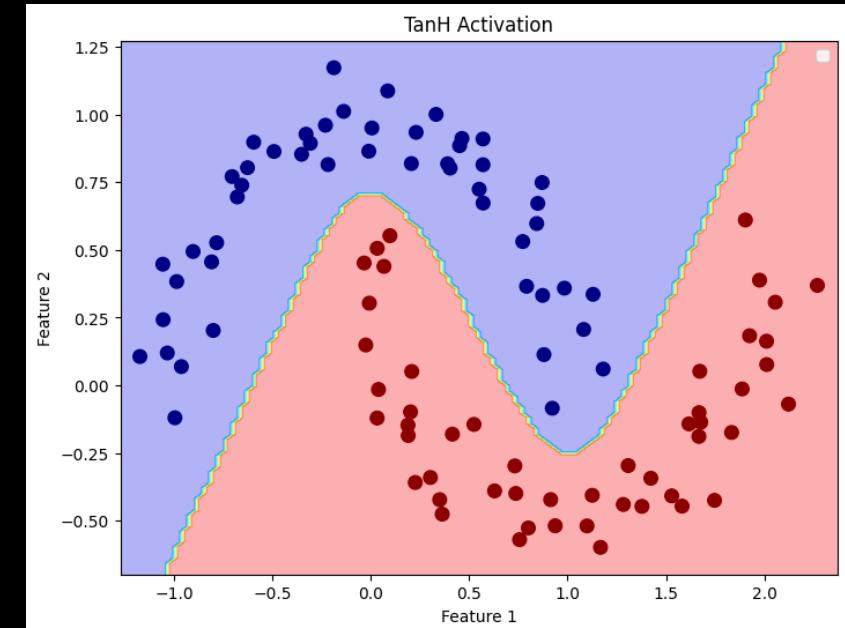
Activation Function introduces **non-linearity** to the network

```
# Define a multi-layer perceptron (MLP)
class MLP(nn.Module):
    def __init__(self):
        super(MLP, self).__init__()
        self.fc1 = nn.Linear(2, 10) # Input layer
        self.fc2 = nn.Linear(10, 1) # Hidden layer
        self.activation = nn.ReLU() # Sigmoid activation function

    def forward(self, x):
        x = self.activation(self.fc1(x))
        x = self.activation(self.fc2(x))
        return x
```

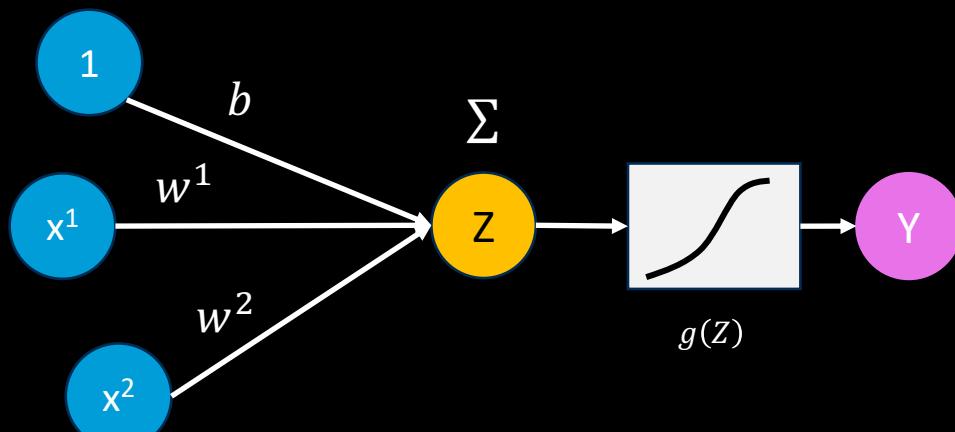


`self.activation = torch.nn.SiLU()`

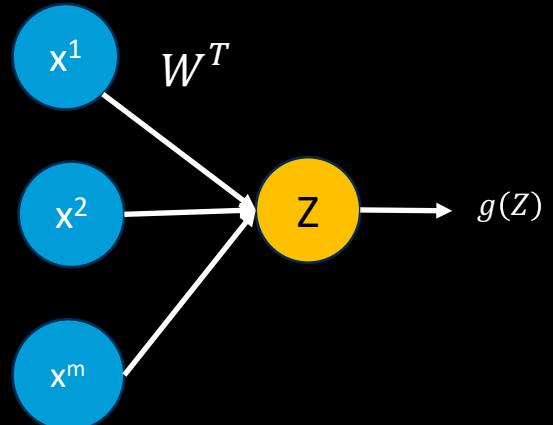


`self.activation = torch.nn.Tanh()`

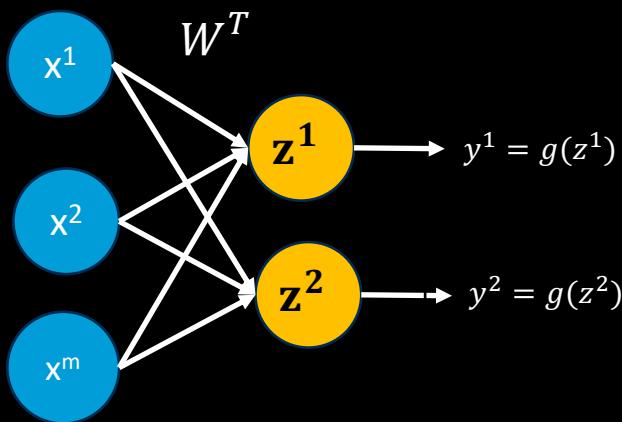
Building Neural Network using Perceptron



Building Neural Network using Perceptron

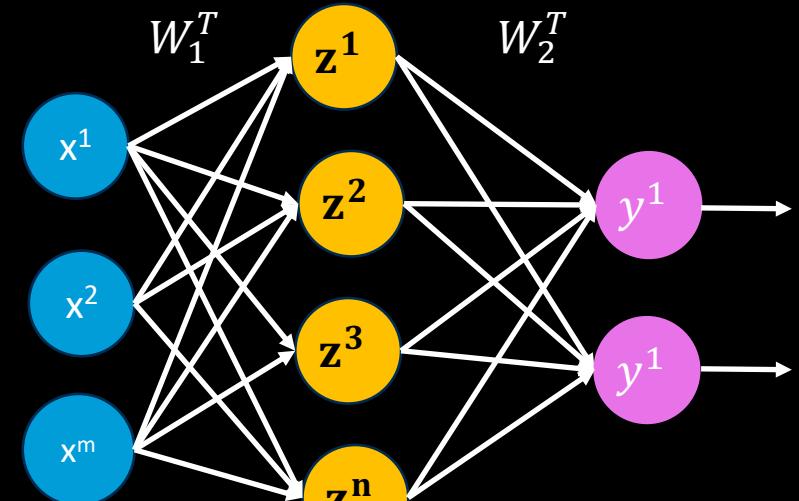


The Perceptron: Simplified



Multi Output Perceptron

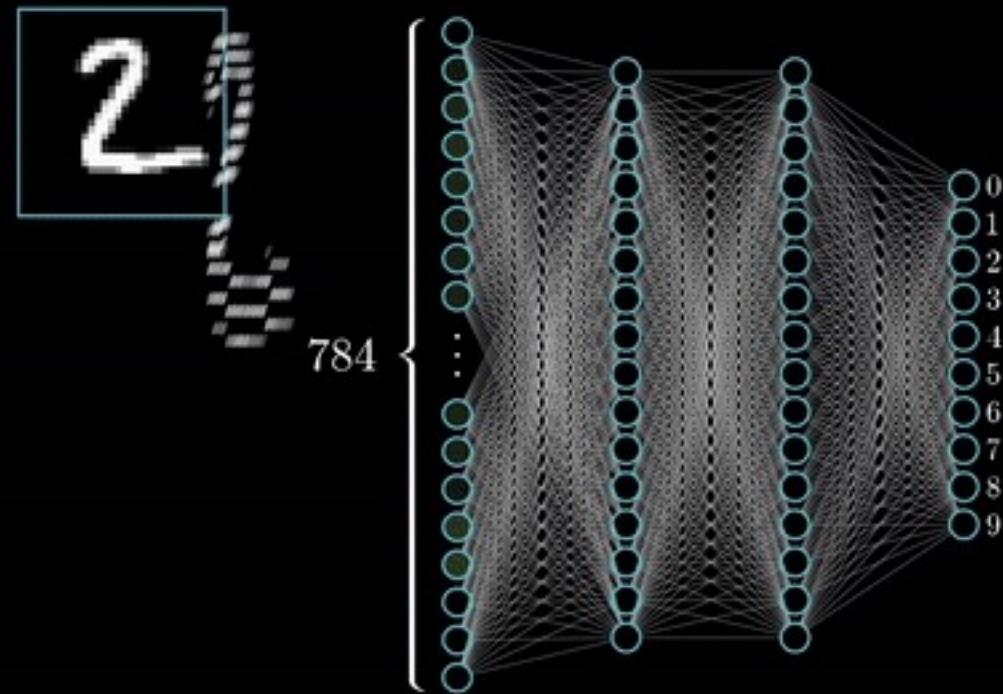
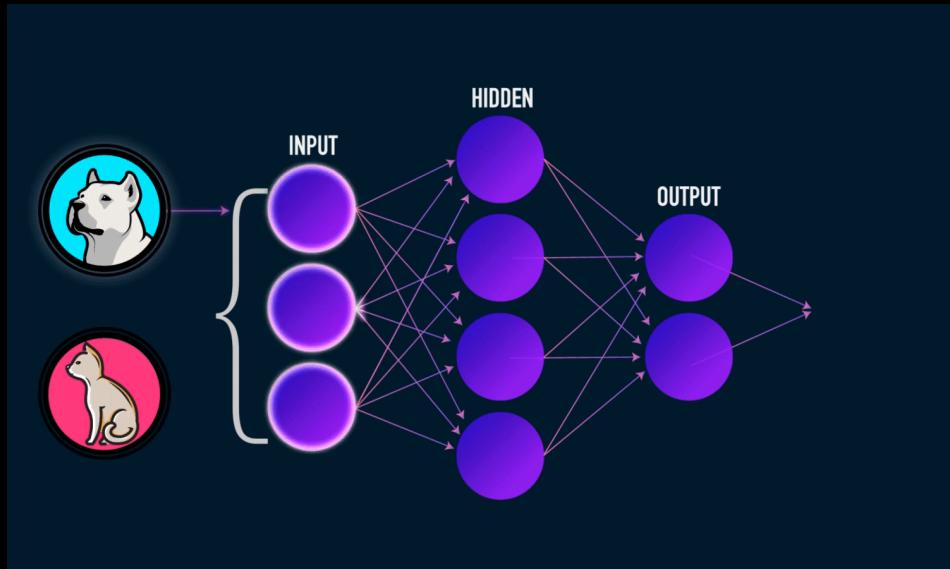
Dense Layers : All inputs are densely connected to all outputs



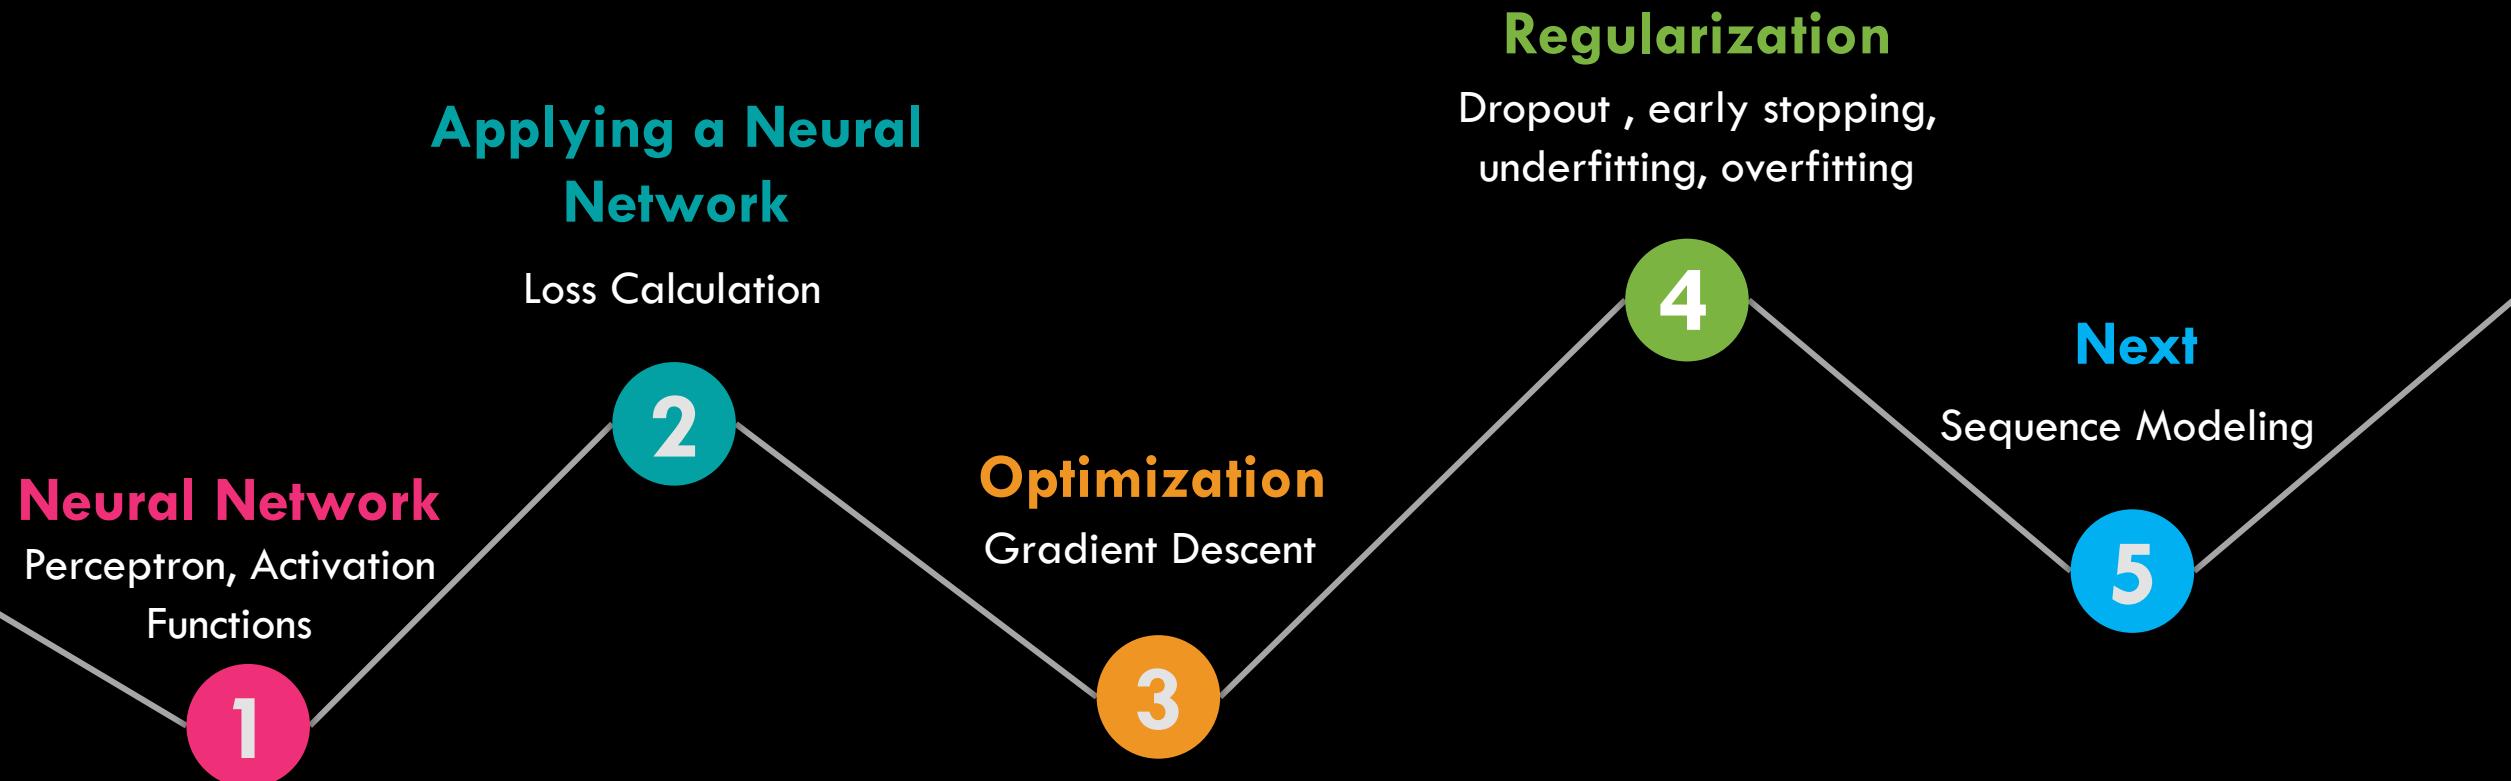
Single Layer Neural Network

Multi-Layer Perceptron

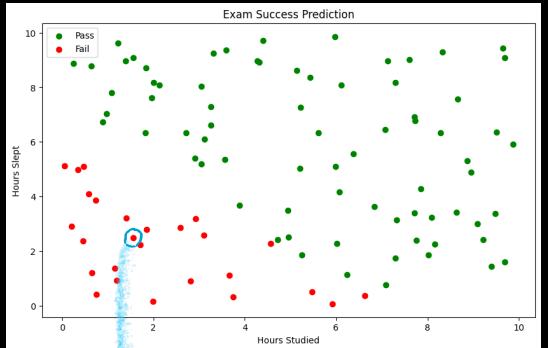
Deep Neural Network



TODAY – Introduction to Deep Learning



Applying a Neural Network



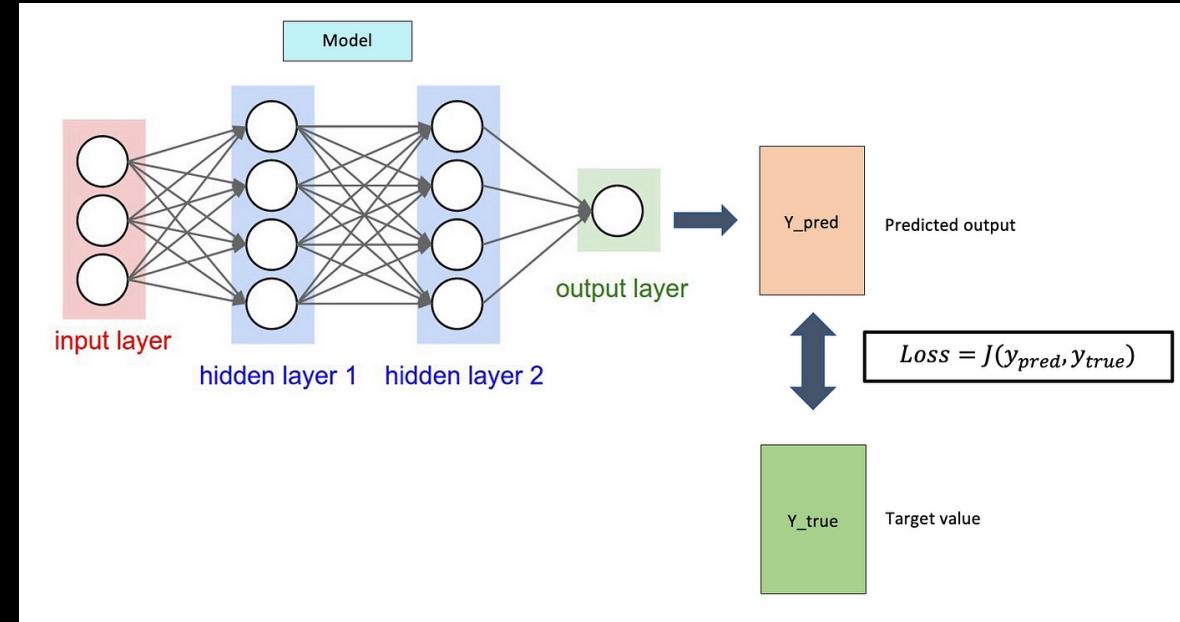
input

$$\begin{bmatrix} \text{studied} \\ \text{slept} \end{bmatrix} = \begin{bmatrix} 1.56 \\ 2.49 \end{bmatrix}$$

Predicted $y = 1.0$

Actual $y = 0.0$

Calculating Loss



$$Loss = J(y_{pred}, y_{true})$$

Binary Cross Entropy Loss

```
criterion = torch.nn.BCEWithLogitsLoss() # Bi  
outputs = model(X_tensor) # Forward pass  
loss = criterion(outputs.squeeze(1), Y_tensor)
```

Mean Squared Error

```
# Define loss function  
criterion = torch.nn.MSELoss()  
# Forward pass  
outputs = model(X_train)  
loss = criterion(outputs, y_train)
```

Mean Absolute Error

```
import torch.nn as nn  
  
mae_loss = nn.L1Loss()
```

- Binary Classification
- Output is probability between 0 and 1

- Regression
- Penalizes larger errors more heavily than small errors

- Regression
- Sensitive to outlier data

Calculating Loss



$$Loss = J(y_{pred}, y_{true})$$

Huber Loss

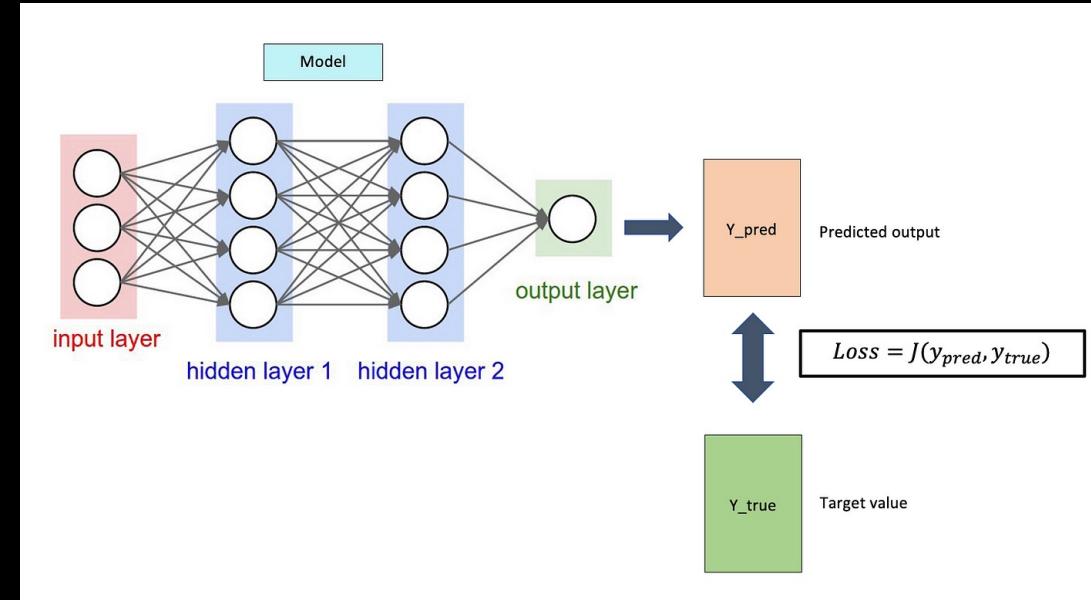
```
import torch.nn as nn  
  
huber_loss = nn.SmoothL1Loss()
```

- Regression
- Combination of L1 and L2 loss

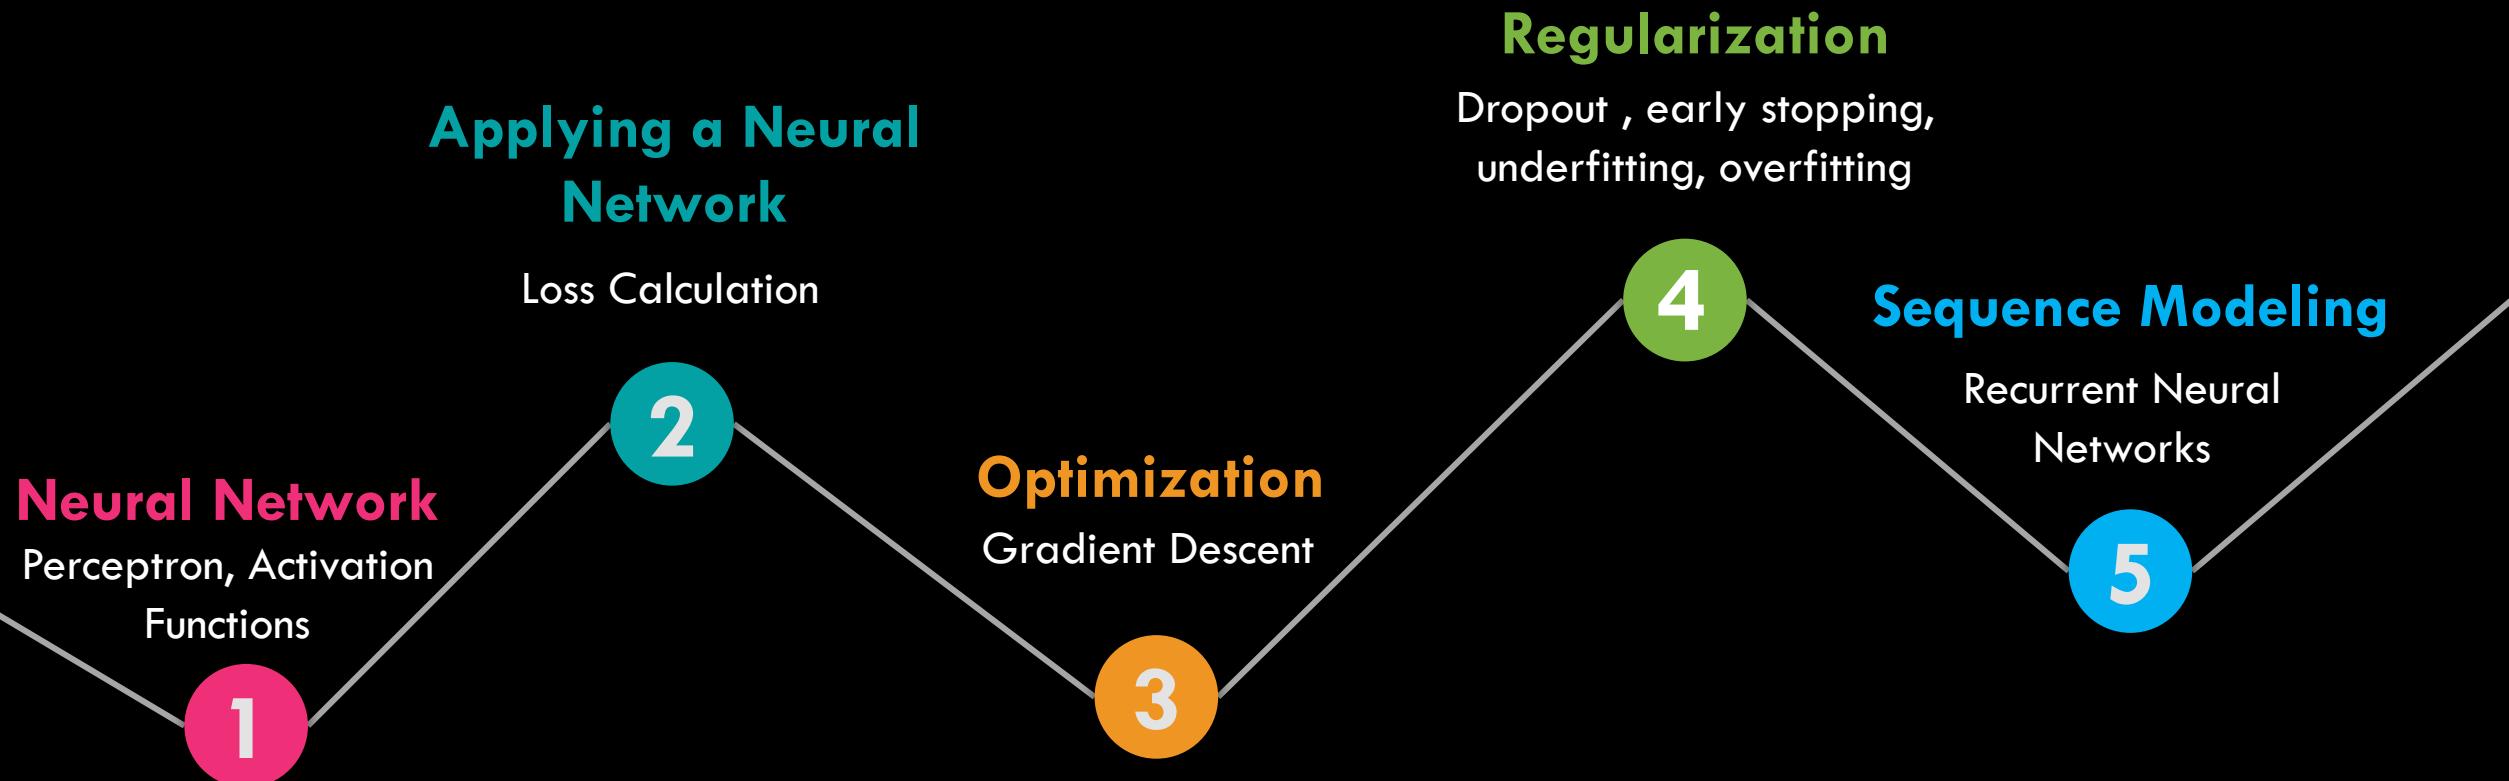
KL Divergence

```
import torch.nn as nn  
  
kl_div_loss = nn.KLDivLoss(reduction='batchmean')
```

- Generative Modeling
- Measures difference between two probability distributions



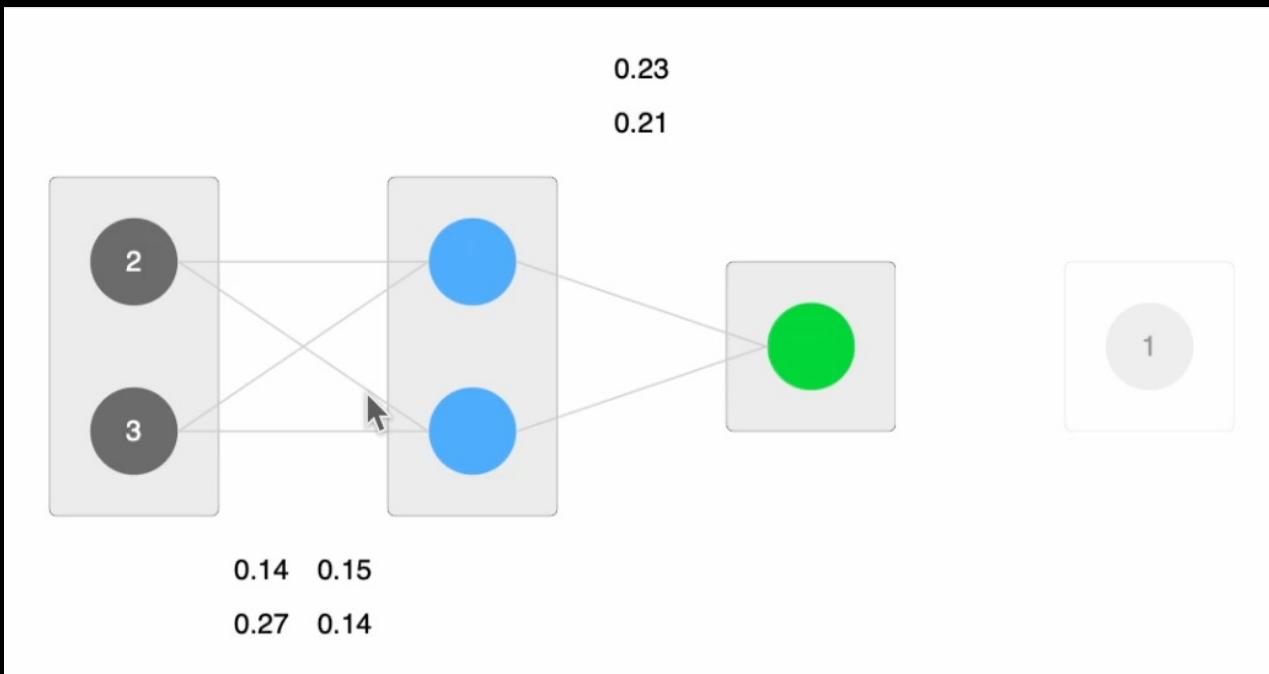
TODAY – Introduction to Deep Learning



Optimization



Forward Pass



Our **OBJECTIVE** is to find new set of weights that decrease the loss.

Loss Optimization

Loss Optimization

Optimization



HoursStudied	HoursSlept	Pass
3.75	0.31	0.00
9.51	6.36	1.00
7.32	3.14	1.00
5.99	5.09	1.00
1.56	9.08	1.00
1.56	2.49	0.00
0.58	4.10	0.00
8.66	7.56	1.00
6.01	2.29	1.00
7.08	0.77	1.00
0.21	2.90	0.00
9.70	1.61	1.00
8.32	9.30	1.00
2.12	8.08	1.00
1.82	6.33	1.00
1.83	8.71	1.00
3.04	8.04	1.00
5.25	1.87	1.00
4.32	8.93	1.00
2.91	5.39	1.00

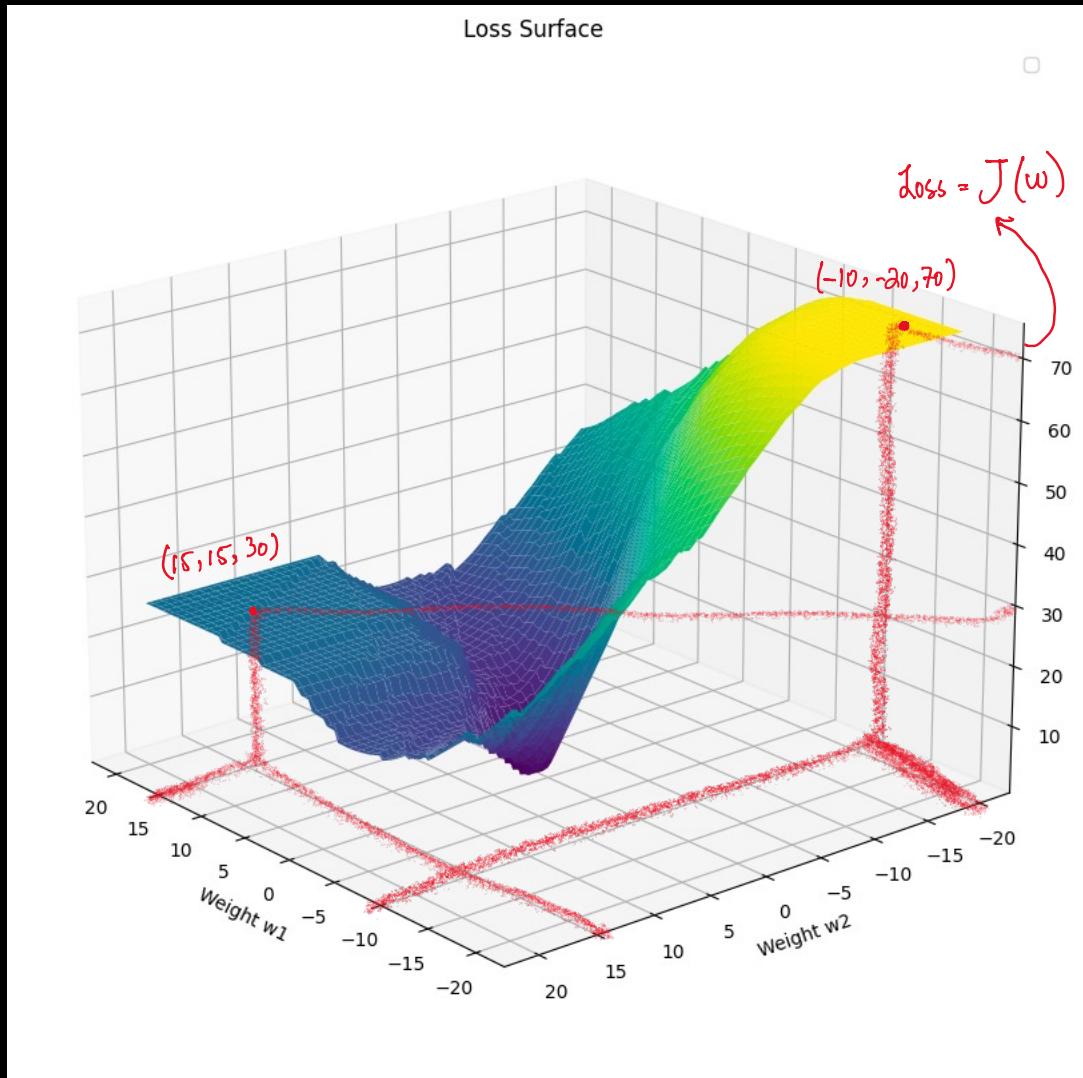
$$\text{Pass} = g(w' \times \text{hours_studied} + w^2 \times \text{hours_slept} + b)$$

```
# Define a mesh grid for w1 and w2
w1_values = np.linspace(-20, 20, 100)
w2_values = np.linspace(-20, 20, 100)
```

Loop → Select w_1, w_2 from above dist.

- Calculate Loss $J(w)$ for all input data

Looping over the
all w_1, w_2 values we get
this 3D plot

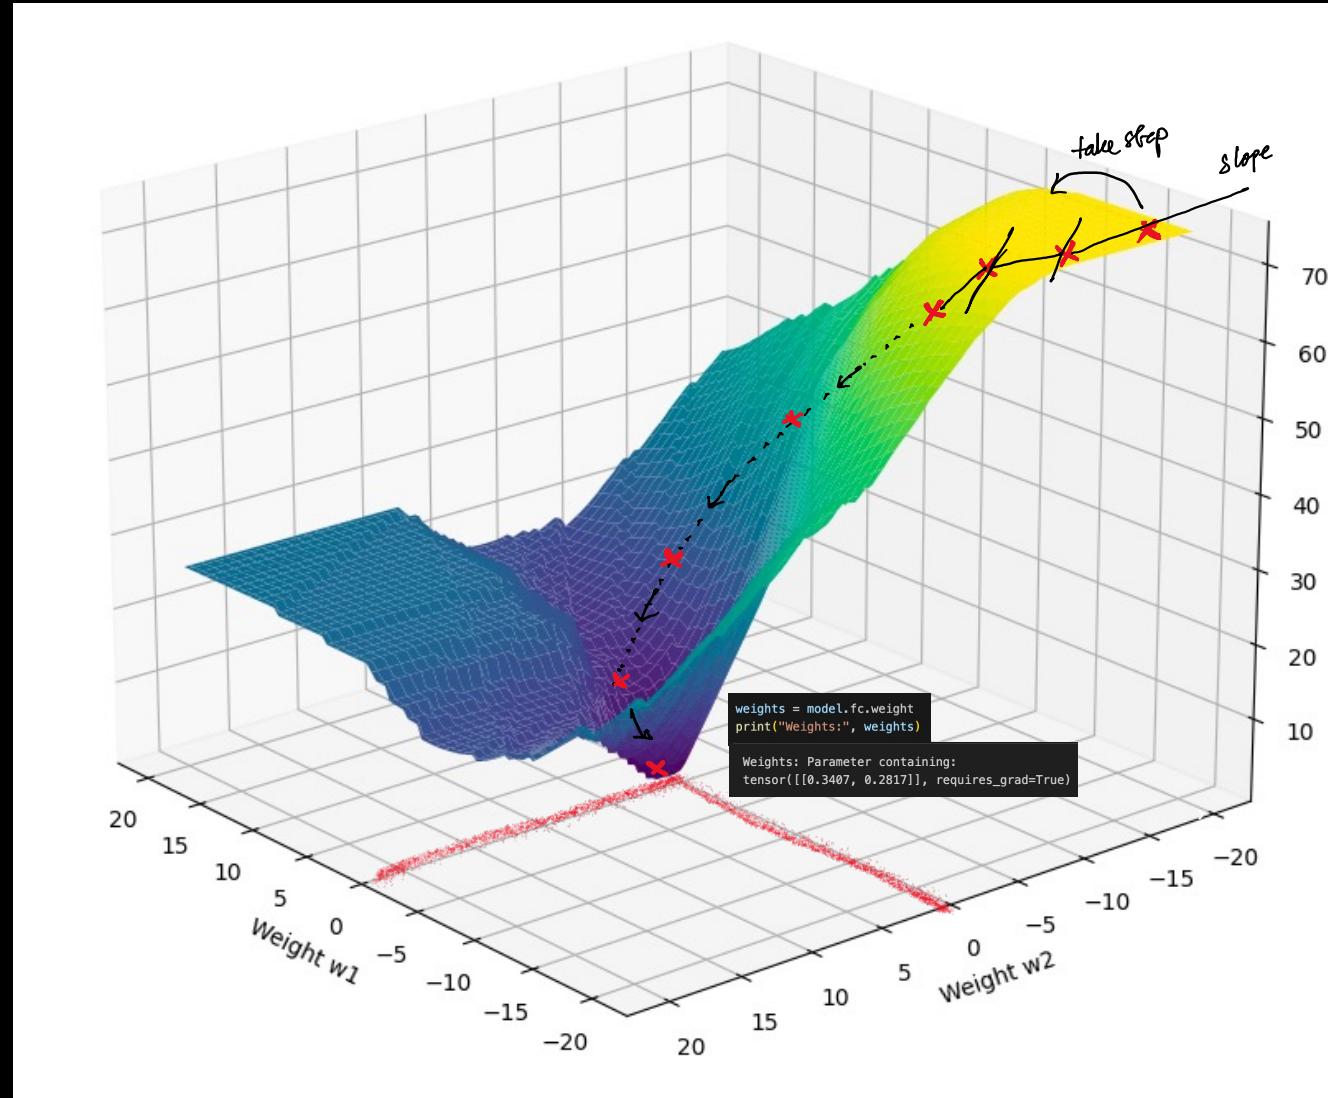


Optimization



- Randomly select value for w_1, w_2
 $w_1 = -20, w_2 = 20$
- Calculate direction of change at this point - gradient
- Take next step opposite direction of slope \rightarrow Repeat above 3 steps

Loss Optimization



Terminology Alert!

These repeated set of actions are Gradient Descent

Step Size is the Learning Rate

Optimization



Gradient Descent

Algorithm

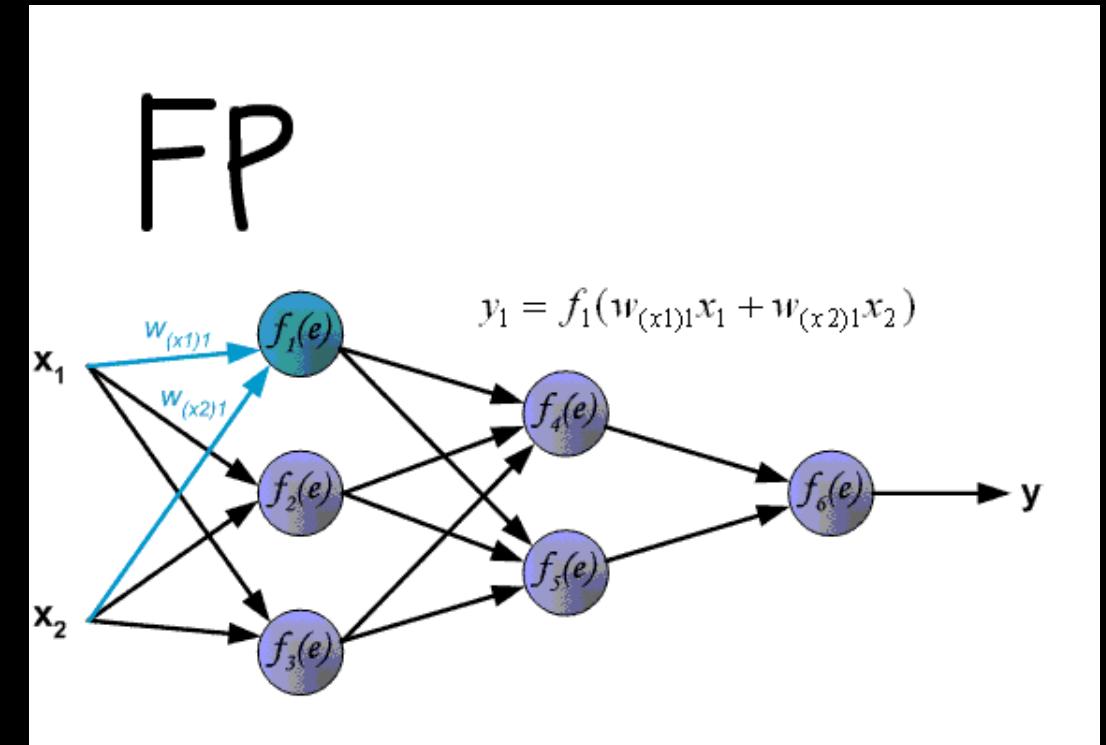
1. Initialize weights randomly

2. Loop until convergence:

3. Compute gradient, $\frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$ *slope*

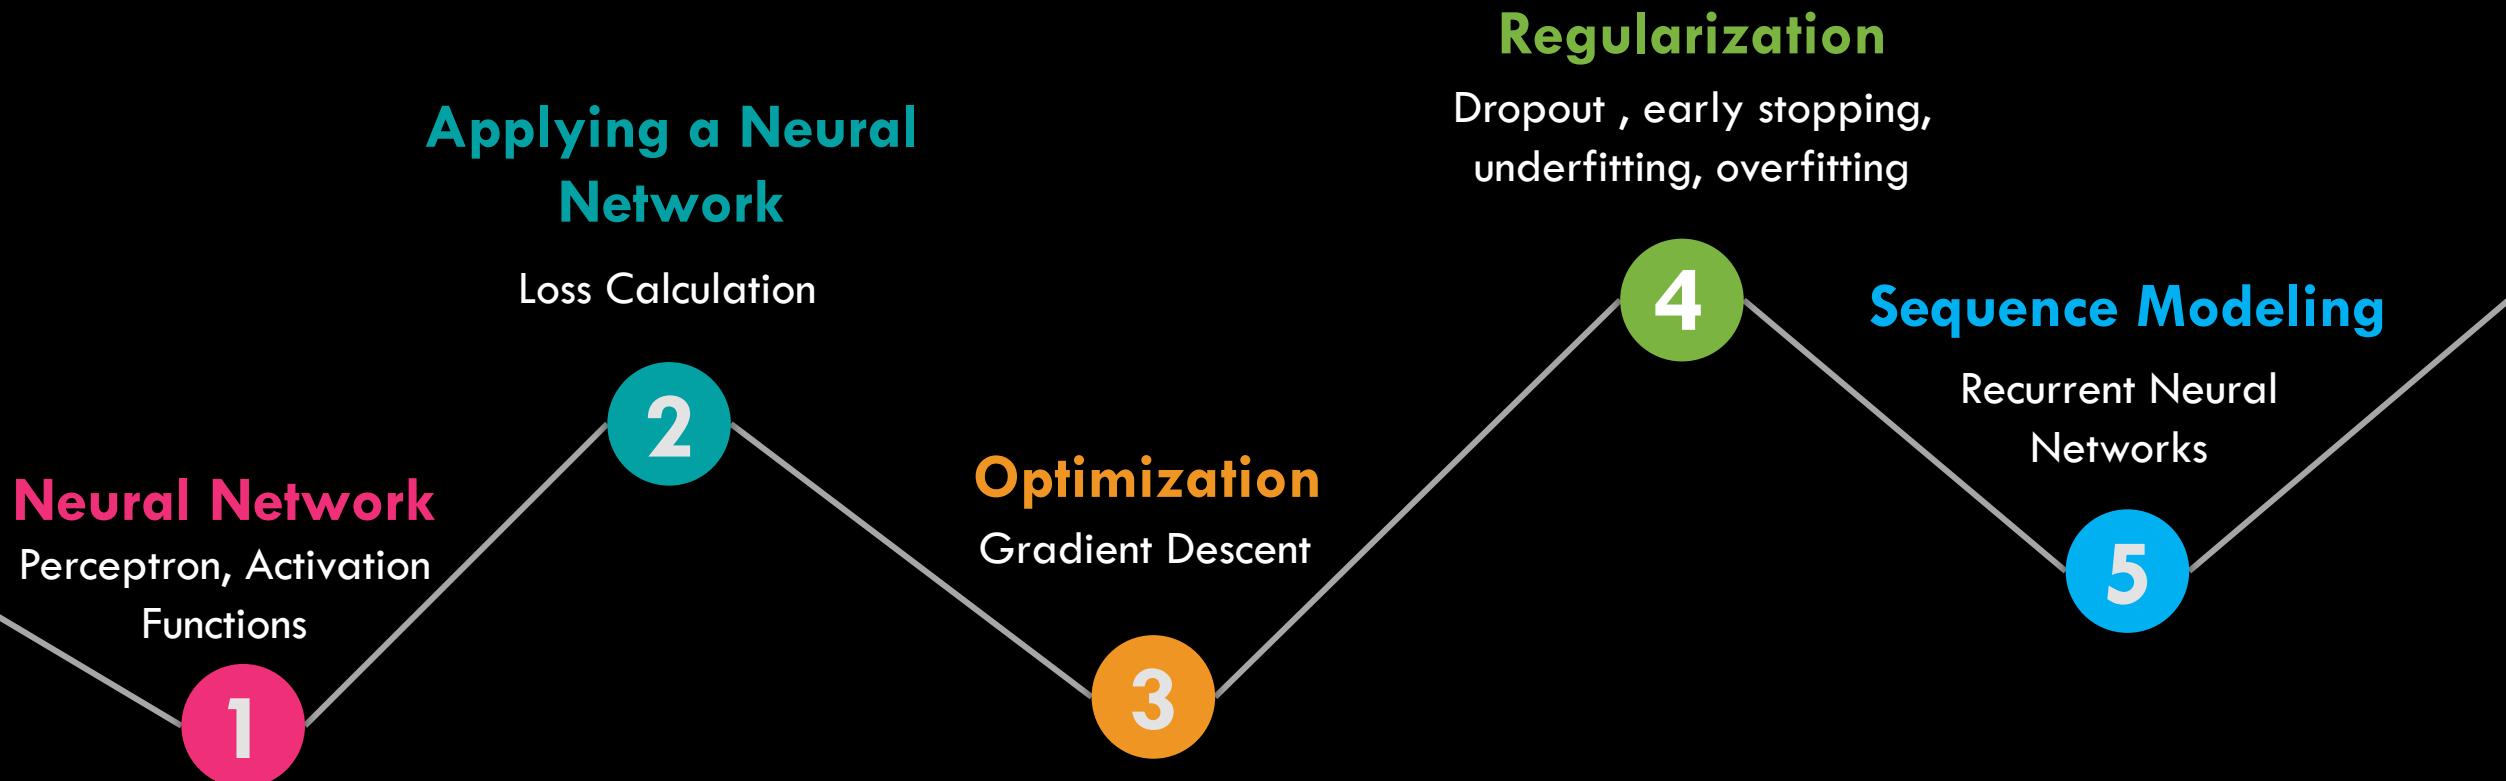
4. Update weights, $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$

5. Return weights



https://miro.medium.com/v2/resize:fit:1144/1*x73oS7hzk6tjv_fhh2IK3g.gif

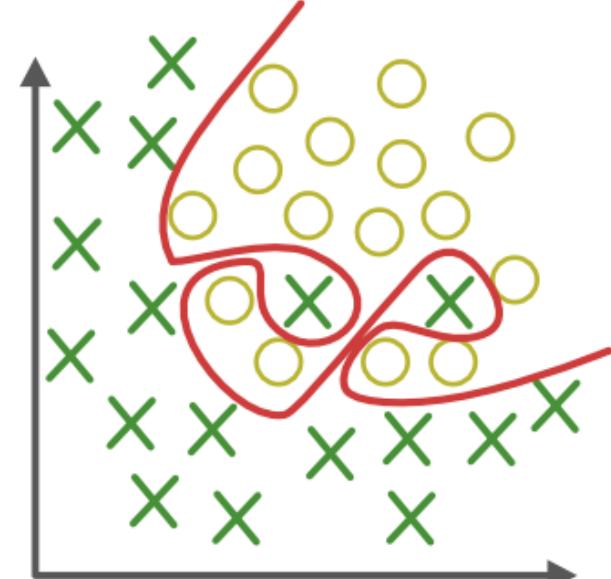
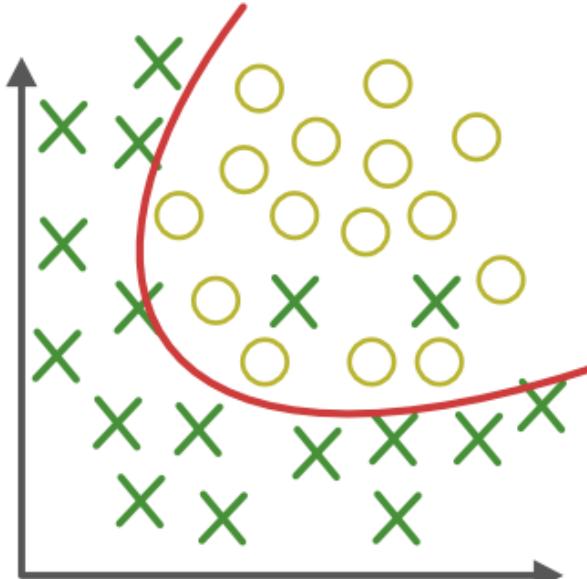
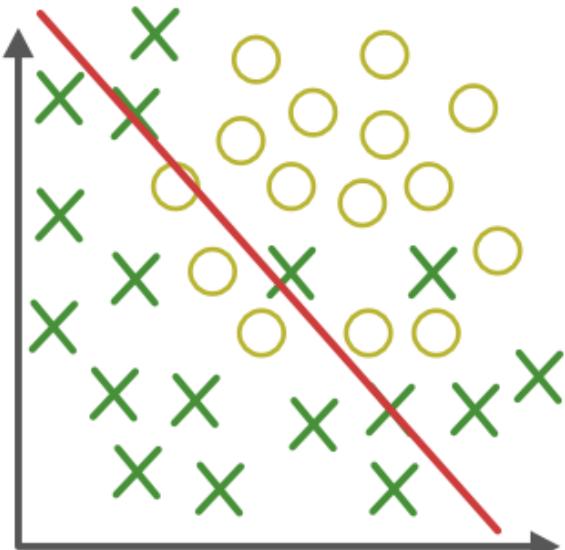
TODAY – Introduction to Deep Learning



Regularization



Dropout , early stopping,
underfitting, overfitting



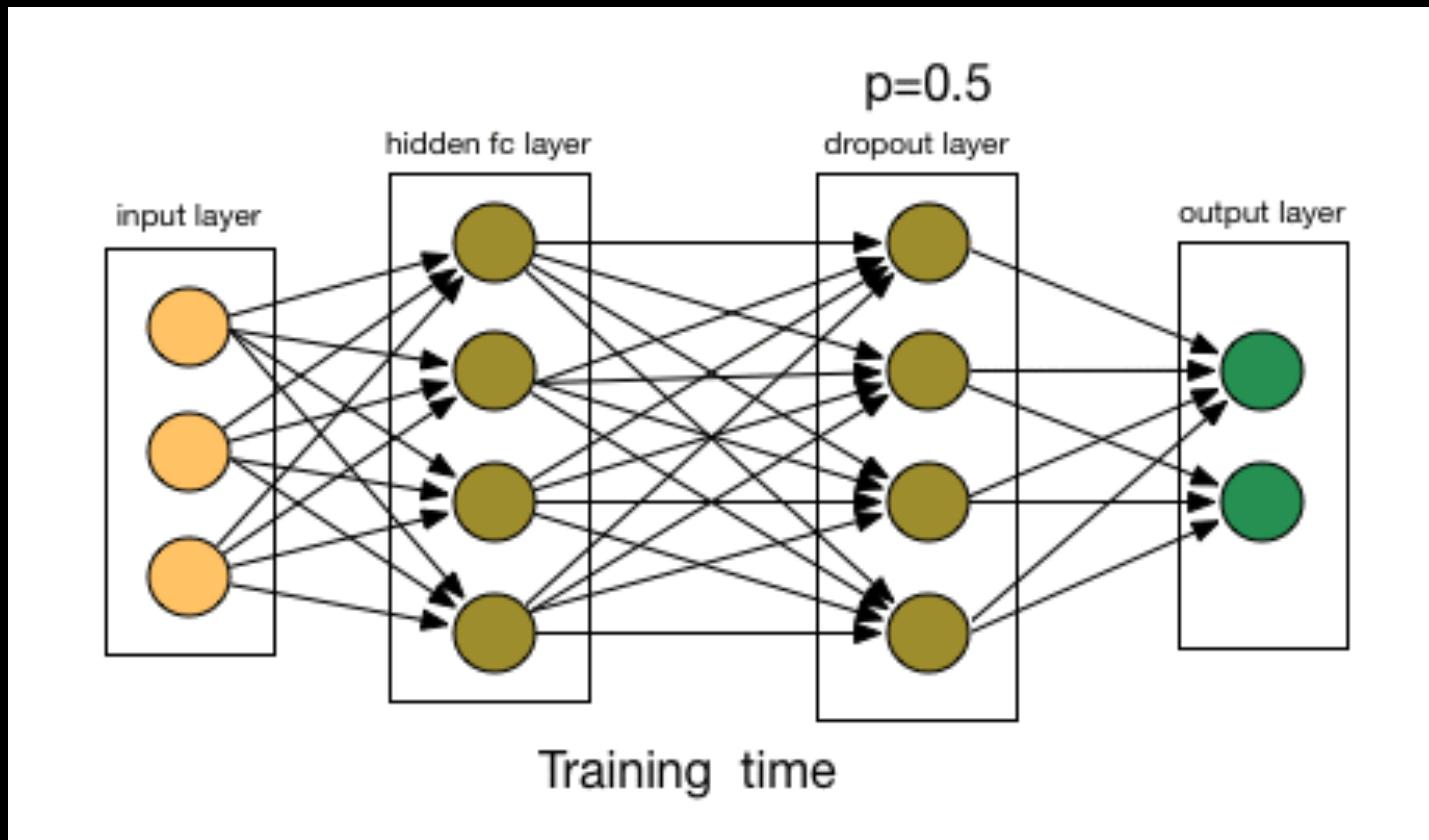
Regularization



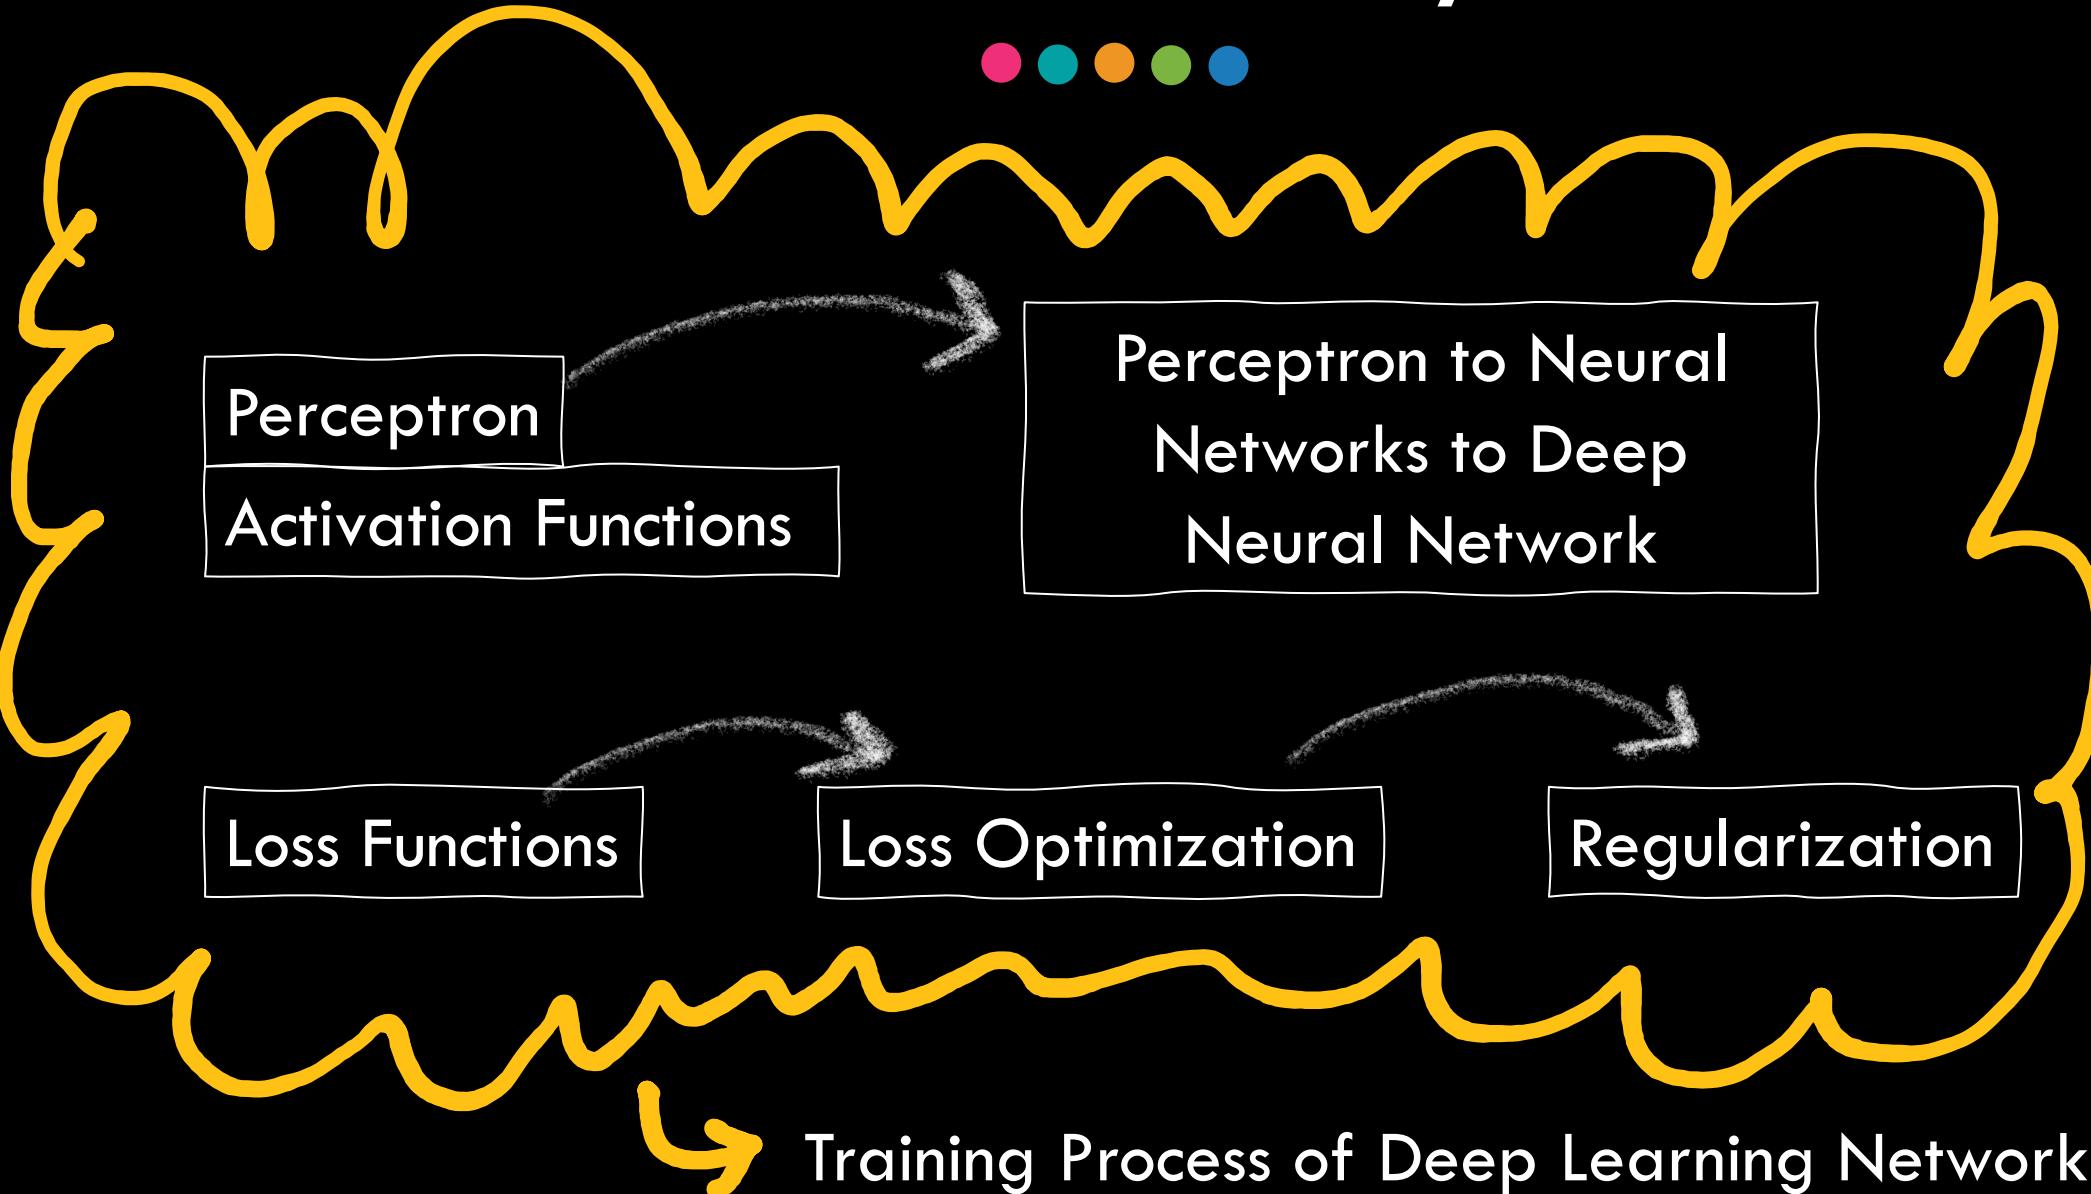
Dropout

During Training randomly set some activations to 0

- Typically drops 50% of activations in a layer
- Forces network to not rely on any one node



Reviewed Today



Putting it all together in Pytorch



```
import torch
import torch.nn as nn
import torch.(module) model_selection, TensorDataset
from torch.utils import TensorDataset
from sklearn.model_selection import train_test_split
import numpy as np

# Convert DataFrame to PyTorch tensors
X = torch.tensor(exam_data[['HoursStudied', 'HoursSlept']].values, dtype=torch.float32)
y = torch.tensor(exam_data['Pass'].values, dtype=torch.float32) → output vector

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Create DataLoader for training and testing sets
train_dataset = TensorDataset(X_train, y_train)
train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
test_dataset = TensorDataset(X_test, y_test)
test_loader = DataLoader(test_dataset, batch_size=32) → pass only 32 input points in one iteration

# Define the perceptron model
class Perceptron(nn.Module):
    def __init__(self, input_size):
        super(Perceptron, self).__init__()
        self.fc = nn.Linear(input_size, 1) # Single output neuron

    def forward(self, x):
        x = torch.sigmoid(self.fc(x)) # Apply sigmoid activation function
        return x

# Initialize the model, loss function, and optimizer
model = Perceptron(input_size=2)
criterion = nn.BCELoss() # Binary cross-entropy loss
optimizer = optim.SGD(model.parameters(), lr=0.01)
```

Prepare Dataset

→ input vector
→ Break the data into train, test
→ pass only 32 input points in one iteration

Define Model

Initialize Model

```
# Training loop
num_epochs = 100
for epoch in range(num_epochs):
    model.train() # Set model to training mode
    for inputs, targets in train_loader:
        optimizer.zero_grad() # Zero the gradients
        outputs = model(inputs) # Forward pass
        loss = criterion(outputs.squeeze(), targets) # Calculate the loss
        loss.backward() # Backward pass
        optimizer.step() # Update weights

    # Evaluate the model on the test set
    model.eval() # Set model to evaluation mode
    with torch.no_grad():
        test_loss = 0
        total_correct = 0
        total_samples = 0
        for inputs, targets in test_loader:
            outputs = model(inputs)
            test_loss += criterion(outputs.squeeze(), targets).item()
            predicted = (outputs >= 0.5).float() # Convert to binary predictions
            total_correct += (predicted == targets).sum().item()
            total_samples += targets.size(0)

    # Print metrics
    test_loss /= len(test_loader)
    accuracy = total_correct / total_samples
    print(f'Epoch {epoch+1}/{num_epochs}, Test Loss: {test_loss:.4f}, Test Acc: {accuracy:.4f}')
```

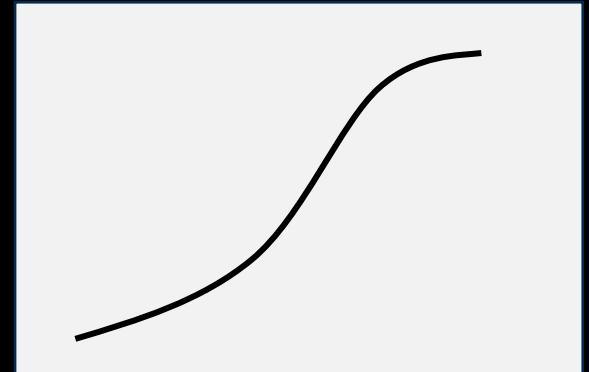
Training Loop

Model Evaluation

References

- <http://introtodeeplearning.com/>
- [lucid by openai](#)
- [microscope by openai](#)
- [Chris Olah's blog](#)

Activation Function



$g(Z)$

